



March 2012

Apache Flume (NG)

Alexander Lorenz | Customer Operations Engineer

cloudera

Overview

- Stream data (events, not files) from clients to sinks
- Clients: files, syslog, avro, ...
- Sinks: HDFS files, HBase, ...
- Configurable reliability levels
 - Best effort: “Fast and loose”
 - Guaranteed delivery: “Deliver no matter what”
- Configurable routing / topology

Architecture

Component	Function
Agent	The JVM running Flume. One per machine. Runs many sources and sinks.
Client	Produces data in the form of events. Runs in a separate thread.
Sink	Receives events from a channel. Runs in a separate thread.
Channel	Connects sources to sinks (like a queue). Implements the reliability semantics.
Event	A single datum; a log record, an avro object, etc. Normally around ~4KB.

Agent

- Runs many clients and sinks
- Java properties-based configuration
- Low overhead (-Xmx20m)
 - But adding RAM increases performance

Sources

- Plugin interface
- Managed by a *SourceRunner* that controls threading and execution model (e.g. polling vs. event-based)
- Included: exec, **avro**, syslog, ...

Sources

```
public class MySource implements PollableSource {
    public Status process() {
        // Do something to create an Event..
        Event e = EventBuilder.withBody(...).build();
        // A channel instance is injected by Flume.
        Transaction tx = channel.getTransaction();
        tx.begin();
        try {
            channel.put(e);
            tx.commit();
        } catch (ChannelException ex) {
            tx.rollback();
            return Status.BACKOFF;
        } finally {
            tx.close();
        }
        return Status.READY;
    }
}
```

Channel

- Plugin interface
- Transactional
- Provide queuing between source / sink
- Provide reliability semantics
 - MemoryChannel: Basically a Java BlockingQueue.
 - JDBC
 - WAL

Sinks

- Plugin interface
- Managed by a *SinkRunner* that controls threading and execution model
- Included: HDFS files (various formats)

Sinks Code

```
public class MySink implements PollableSink {
    public Status process() {
        Transaction tx = channel.getTransaction();
        tx.begin();
        try {
            Event e = channel.take();
            if (e != null) {
                // ...
                tx.commit();
            } else {
                return Status.BACKOFF;
            }
        } catch (ChannelException ex) {
            tx.rollback();
            return Status.BACKOFF;
        } finally {
            tx.close();
        }
        return Status.READY;
    }
}
```

Tiered collection

- Send events from agents to another tier of agents to aggregate
- Use an Avro sink (really just a client) to send events to an Avro source (really just a server) in another machine
- Failover supported
- Load balancing (soon)
- Transactions guarantee handoff

Tiered collection – The handoff

- Agent 1: Tx begin
- Agent 1: Channel take event
- Agent 1: Sink send
- Agent 2: Tx begin
- Agent 2: Channel put
- Agent 2: Tx commit, respond OK
- Agent 1: Tx commit (or rollback)

Configuration

- done in a single file
- identifier for flows
- `<identifier>.type.subtype.parameter.config`, where `<identifier>` is the name of the agent
- flow has a client, type, channel, sink
- can have multiple channels in one flow

Simple Configuration Example

```
syslog-agent.sources = Syslog
syslog-agent.channels = MemoryChannel-1
syslog-agent.sinks = Console

syslog-agent.sources.Syslog.type = syslogTcp
syslog-agent.sources.Syslog.port = 5140

syslog-agent.sources.Syslog.channels =
MemoryChannel-1
syslog-agent.sinks.Console.channel = MemoryChannel-1

syslog-agent.sinks.Console.type = logger
syslog-agent.channels.MemoryChannel-1.type = memory
```

HDFS Configuration Example

```
syslog-agent.sources = Syslog
syslog-agent.channels = MemoryChannel-1
syslog-agent.sinks = HDFS-LAB

syslog-agent.sources.Syslog.type = syslogTcp
syslog-agent.sources.Syslog.port = 5140

syslog-agent.sources.Syslog.channels = MemoryChannel-1
syslog-agent.sinks.HDFS-LAB.channel = MemoryChannel-1

syslog-agent.sinks.HDFS-LAB.type = hdfs

syslog-agent.sinks.HDFS-LAB.hdfs.path = hdfs://NN.URI:PORT/
flumetest/'%{host}''
syslog-agent.sinks.HDFS-LAB.hdfs.file.Prefix = syslogfiles
syslog-agent.sinks.HDFS-LAB.hdfs.file.rollInterval = 60
syslog-agent.sinks.HDFS-LAB.hdfs.file.Type = SequenceFile
syslog-agent.channels.MemoryChannel-1.type = memory
```

Features

- Fan out: One source, many channels
- Fan in: Many sources, one channel
- Processors (aka: decorators)
- Auto-batching of events in RPCs...
- Multiplexing channels for datamining
- Avro implementation in both ways

Thank You

- **Web:** <https://cwiki.apache.org/FLUME/getting-started.html>
- **ML:** flume-user@incubator.apache.org
- **Mail:** alexander@cloudera.com
- **Blog:** mapredit.blogspot.com
- **Twitter:** @mapredit